

My eval corpus had 50 traces and zero spans.

The cause was not what I expected.

WHERE I WAS LOOKING

workspace/<run_id>/

src/, tests/, and no telemetry of any kind.

440 directories. 221 of them completely empty.

This is where my corpus builder was pointed.

WHERE THE DATA WAS

output/runs/<run_id>/

211 dated run directories. 210 run.json carrying backend, completion time and final status.

141 state.json with real cost and token counts. 60 cost logs. 8 receipts.

Three backends. A 69-day span. Untouched since July.

THE WHOLE CAUSE

One default argument.

```
evals/cli.py:547  
backfill.add_argument(  
    "--workspace-root",  
    default="./workspace")
```

workspace/ is where the agents' generated code lands.

output/runs/ is where the harness writes its own records.

I was reading the first one.

WHY IT HID

Nothing fails loudly.

A writer and a reader that disagree about where data lives do not raise an error.

The corpus simply comes back thinner than the system that produced it — and every statistic above it inherits that, quietly.

Mine looked like 50 rows. It was closer to $n=1$.

Does the thing that writes your traces agree with the thing that reads them?

Ten minutes to check. It is now the first question I ask of any eval setup, mine included.

Full write-up, with the queries, in this week's issue.

I write up one measurement from a real multi-agent system every week, with the query that produced it.