

Five questions before you trust a green check.

Each has a cheap answer and an expensive one. The cheap answer is usually wrong.

00

Does the thing that writes your traces agree with the thing that reads them?

Two directories, one writer, one reader, and no test between them.

Nothing fails loudly when they diverge.

01

How many real traces, and how diverse?

Group by every dimension you have — backend, scenario, status, date.

If one value dominates a column, your n is a fiction.

02

How many has a human actually read?

Not scored. Read, with notes.

This is the number no dashboard reports and no survey asks for.

03

Where did your failure categories come from?

If you cannot point from a category to the specific traces that produced it, you are testing your imagination.

04

What does a green run actually assert? Say it in one sentence.

"My check code behaves as written" is real and useful.

It is not a quality claim. Stop letting it travel as one.

I failed all five on a system I was proud of.

13,000 lines of eval infrastructure, built before I asked what it was measuring.

The write-up has the queries, the numbers, and what I changed.

I write up one measurement from a real multi-agent system every week, with the query that produced it.